

Intelligent Document Processing (IDP) - 12 Hour Hackathon Roadmap

Team Size: 4 Members

Hackathon Start Time: 11 PM

1. Overview

This document provides a complete 12-hour execution roadmap, task splitting, technology stack, and presentation preparation strategy for building an Intelligent Document Processing system during an overnight hackathon. The goal is to build a lean and functional MVP that supports:

- Document upload (PDF, DOCX, Images)
 - Text extraction
 - Summarization
 - Document classification
 - Automatic routing
 - Semantic search using FAISS
 - Simple, impressive UI (Streamlit)
-

2. Technology Stack

Backend:

- Python
- FastAPI
- LangChain
- FAISS Vector Store
- PyPDF2 / pdfplumber (PDF extraction)
- docx2txt (DOCX extraction)

- Optional: Tesseract or EasyOCR for image OCR

AI/ML:

- Gemini 2.0 Flash / Pro (Summarization & Classification)
- Google Generative AI Embeddings
- SentenceTransformers
- Optional ML classifier (RandomForest / Logistic Regression)

Frontend:

- Streamlit (recommended for hackathon speed)
- HTML/CSS optional for enhancements

Storage:

- Local storage for uploaded docs
- FAISS index locally stored as faiss_index.bin
- Optional SQLite DB for metadata

3. Team Division (4 Members)

Member 1 – Backend & Extraction

- Build FastAPI endpoints: /upload, /extract
- Implement PDF, DOCX extraction logic
- Optional: Integrate OCR using easyocr
- Implement text cleaning pipeline

Member 2 – AI Pipeline (Summarization, Classification, Embeddings)

- Build summarization function using Gemini
- Build classification prompt for document types
- Implement routing logic (rule-based + LLM)

- Create chunking + embedding + FAISS indexing
- Build /summarize, /classify, /search endpoints

Member 3 – UI/Frontend Developer

- Build Streamlit UI
- Document upload interface
- Display extracted text, summary, category, routing
- Build semantic search UI
- Add document history visual panel

Member 4 – Integrations, QA, Pitch Deck

- Integrate Streamlit with FastAPI
 - Build demo script
 - Create slides (Problem → Solution → Demo → Impact)
 - Prepare sample documents
 - Conduct testing and optimizations
-

4. 12-Hour Smart Execution Timeline

11 PM – 12 AM (1 hour) — Planning & Setup

- Finalize features of the MVP
 - Create GitHub repo
 - Set up FastAPI skeleton
 - Set up Streamlit skeleton
 - Confirm folder structure
 - Decide classification categories
-

12 AM – 3 AM (3 hours) — Backend + AI Foundation

Member 1:

- Implement /upload
- Implement text extraction for PDF & DOCX
- Test extraction on 3 sample files

Member 2:

- Implement chunking + embedding pipeline
 - Build FAISS index creation & update
 - Write summarization prompt
 - Write classification prompt
-

3 AM – 6 AM (3 hours) — Frontend + API Integration

Member 3:

- Build Streamlit UI components:
- Upload section
- Display extracted text
- Display summary
- Display routing & category
- Search page

Member 4:

- Integrate UI ↔ backend endpoints
 - Build search workflow
 - Fix request/response formatting
 - Verify full flow: Upload → Extract → Summarize → Classify → Search
-

6 AM – 8 AM (2 hours) — Polish & Additional Features

All Members:

- Add similarity search between documents
 - Improve UI layout
 - Add error handling & loading states
 - Add multi-document search sorting
 - Improve prompts for classification
-

8 AM – 10 AM (2 hours) — Testing, Sample Docs, Pitch Deck

Member 4:

- Build pitch slides:
 1. Problem Statement
 2. Why it matters
 3. Proposed Solution
 4. Architecture Diagram
 5. Demo Flow
 6. Impact & Future Scope

Members 1–3:

- Test all features with:
 - HR letter
 - Invoice
 - Engineering doc
 - Purchase order
 - Safety document
 - Final fixes

10 AM – 11 AM (1 hour) — Final Dry Run

- Run the whole demo end-to-end
 - Ensure UI looks clean
 - Fix performance & formatting issues
 - Finalize GitHub README
 - Prepare for judging
-

5. Final Deliverables for Judges

1. Working demo:

- Document upload → extraction → summary → classification → routing → semantic search

2. Clean UI/UX using Streamlit

3. Clear architecture:

- FastAPI Backend
- FAISS Vector Search
- Gemini AI for NLP

4. Pitch Deck with impact and future scope

5. GitHub Repository with:

- Code
- Instructions
- Sample docs
- Demo screenshots

6. Conclusion

This roadmap ensures that a team of 4 can confidently build an impressive, fully functional Intelligent Document Processing MVP within a 12-hour hackathon window. Proper division of tasks, simplified architecture, and leveraging LLMs allow rapid development while maintaining high impact.